

VIKUNA TECHNOLOGIES

Infrastructure & Operations Documentation

Version 3.0 | May 2026 | FINAL

Hostinger KVM2 (×2) — India Data Center

Docker + PostgreSQL 17 + Qwen3 4B Self-Hosted LLM

Confidential — Vikuna Technologies

1. Architecture Overview

Vikuna infrastructure runs on two Hostinger KVM2 VPS instances in the India Data Center, orchestrated via Docker Compose. The two servers have distinct roles — one handles product databases and APIs, the other runs the shared LLM inference service.

1.1 Two-VPS Architecture

VPS	IP Address	Hostname	Primary Role	Expires
Main VPS	187.127.136.65	srv1528480.hstgr.cloud	PostgreSQL + PostgREST + Nginx + DristiQ	Mar 2027
LLM VPS	72.60.222.136	srv1096269.hstgr.cloud	Qwen3 LLM + N8N + Traefik	Nov 2026

1.2 Products Hosted

Product	Database	Prefix	Size	VPS	Status
DristiQ (KaalaDristi)	kaala_dristi_db	KD_	619 MB / 5.6M rows	Main	LIVE ✓
KI-Prime	ki_prime_db	KI_	Empty	Main	Ready
FamilyKnows	familyknows_db	FK_	Empty	Main	Ready
VaNi Framework	vani_db	VN_	Empty	Main	Ready
N8N Workflows	n8ndb (PG 15)	—	~50 MB	LLM	LIVE ✓

1.3 Server Specifications (Both VPS)

Attribute	Value
Provider	Hostinger KVM2
OS	Ubuntu 24.04.4 LTS
RAM	8 GB
Disk	96 GB NVMe
Data Center	India
PostgreSQL	17.9 (Main VPS) / 15 (LLM VPS — N8N only)
LLM	Qwen3 4B Q4_K_M (2.5 GB GGUF)
Cost	₹799/mo each ₹1,598/mo total

2. Main VPS — 187.127.136.65

2.1 Running Containers

Container	Image	Port	RAM Est.	Status
vikuna-postgres	postgres:17-alpine	5432	3–4 GB	Running ✓
vikuna-postgrest	postgrest/postgrest:v12.2.3	3000 (internal)	200 MB	Running ✓
vikuna-nginx	nginx:alpine	80, 443	100 MB	Running ✓
kd-pipeline-api2	python/fastapi (custom)	8101	500 MB	Running ✓

2.2 Pending Services

Service	Purpose	When
node-backend	KI-Prime + FamilyKnows + VaNi APIs	When products ready
pg-backup	Nightly backup to Google Drive	Next session

2.3 API Routing (Nginx)

URL Path	Backend	Purpose
/health	Nginx direct	Health check
/db/*	PostgREST :3000	Database REST API
/api/kd/*	FastAPI :8101	DristiQ compute + panchang APIs
/api/ki/*	Node :4000 (pending)	KI-Prime APIs
/api/fk/*	Node :4000 (pending)	FamilyKnows APIs
/api/vani/*	Node :4000 (pending)	VaNi Framework APIs

2.4 Firewall Rules

Port	Access	Purpose
22/tcp	Everyone	SSH
80/tcp	Everyone	Nginx HTTP
443/tcp	Everyone	Nginx HTTPS (future)
5432/tcp	49.37.150.141 + IPv6 subnet	PostgreSQL direct (Beekeeper)
8080/tcp	Internal only	LLM (removed — now on LLM VPS)

2.5 File Locations

Path	Contents
/opt/vikuna/docker/docker/	Docker compose + configs
/opt/vikuna/docker/docker/.env	All passwords and secrets
/opt/vikuna/docker/docker/config/	Nginx, PG, backup configs

Path	Contents
/opt/vikuna/docker/scripts/	Management + migration scripts
/opt/vikuna/apps/frontends/	React build outputs

2.6 Database Migration Status

Migrations M001–M070 executed on kaala_dristi_db. Key tables:

Table	Purpose	Status
km_astro_rule_master	Canonical astro rules	Live ✓
km_astro_calendar_2026	Full 2026 panchang events	Live ✓
km_astro_daily_signal	Daily computed signals	Live ✓
km_planetary_positions	Ephemeris (Lahiri / Ujjain)	Live ✓
km_finastro_alerts	Finastro alert layer	Pending
km_finastro_muhurta	Muhurta selection	Pending
km_astro_correlation	Correlation engine results	Pending

NOTE: CLAUDE.md is outdated — migrations M063–M070 undocumented. Update before next sprint.

3. LLM VPS — 72.60.222.136

3.1 Running Containers

Container	Image	Port	RAM Est.	Status
vikuna-llm	ghcr.io/ggml-org/llama.cpp:server	8080	~1.6 GB	Running ✓
root-n8n-1	docker.n8n.io/n8nio/n8n	5678 (local)	300 MB	Running ✓
root-n8n-worker-1/2/3	n8nio/n8n	5678	600 MB	Running ✓
n8n-postgres	postgres:15	5432 (internal)	500 MB	Running ✓
redis	redis:6	6379 (internal)	100 MB	Running ✓
root-traefik-1	traefik	80, 443	100 MB	Running ✓

3.2 LLM Configuration

Attribute	Value
Model	Qwen3-4B-Instruct
Quantization	Q4_K_M (2.5 GB file)
Framework	llama.cpp server

Attribute	Value
Context Size	4096 tokens
Threads	3
Parallel Slots	4 (auto)
RAM Usage	~1,575 MiB projected
RAM Available	~6.1 GB free (8 GB total)
Thinking Mode	Disabled via /no_think in system prompt
Model File	/opt/vikuna/models/Qwen3-4B-Q4_K_M.gguf
API Format	OpenAI-compatible
Public Endpoint	https://llm.dristiq.io/v1/chat/completions
API Key	vk-llm-d0efccfd15c0b8fd72214d0b9182032f84106b43

NOTE: Always include /no_think in the system prompt for DristiQ and client project calls to suppress chain-of-thought tokens in structured outputs.

3.3 Traefik Configuration

Route	Target	SSL
llm.dristiq.io	vikuna-llm :8080	Auto (mytischallenge)
n8n.srv1096269.hstgr.cloud	root-n8n-1 :5678	Auto (mytischallenge)

3.4 File Locations

Path	Contents
/root/docker-compose.yml	Main compose file (all services)
/root/.env	Domain, SSL email, credentials
/opt/vikuna/models/	Qwen3 GGUF model file (2.5 GB)

4. Connection Details

NOTE: Store these credentials securely. Never commit to Git.

4.1 Main VPS — PostgreSQL

Field	Value
Host	187.127.136.65
Port	5432
Admin User	vikuna_admin
Admin Password	Vikuna2026Secure

Field	Value
KD App User	kd_app
KD App Password	KdApp2026Secure
Connection String	postgres://vikuna_admin:Vikuna2026Secure@187.127.136.65:5432/<db>

4.2 Main VPS — PostgREST

Field	Value
Base URL	http://187.127.136.65/db/
Health	http://187.127.136.65/health
JWT Secret	21RkWGQnQ0Guwj8XI6j3U+jSgAw...WQ==

4.3 LLM VPS — Qwen3 API

Field	Value
Public Endpoint	https://llm.dristiq.io/v1/chat/completions
Local Endpoint	http://localhost:8080/v1/chat/completions
API Key	vk-llm-d0efccfd15c0b8fd72214d0b9182032f84106b43
Auth Header	Authorization: Bearer vk-llm-d0efccfd15c0b8fd72214d0b9182032f84106b43
Model String	Use any string — llama.cpp ignores it

4.4 LLM VPS — N8N

Field	Value
URL	https://n8n.srv1096269.hstgr.cloud
DB User	n8n
DB Password	xyXdPYrIA5UQGnwK4zJN
Encryption Key	td1z5lJQzDcSQTxwsiAv3GlcBgD0jaok
N8N local LLM	http://vikuna-llm:8080 (same Docker network)

4.5 Beekeeper Studio (Main VPS)

Field	Value
Host	187.127.136.65 (direct)
Port	5432
User	vikuna_admin
Password	Vikuna2026Secure
SSL	Off

5. Commands Reference

5.1 Main VPS — Docker

Command	Description
<code>docker compose up -d</code>	Start all services
<code>docker compose down</code>	Stop all services
<code>docker compose ps</code>	Show status
<code>docker compose restart <service></code>	Restart a service
<code>docker compose logs -f <service></code>	Follow logs
<code>docker stats --no-stream</code>	RAM/CPU snapshot

5.2 Main VPS — PostgreSQL

Command	Description
<code>docker exec vikuna-postgres psql -U vikuna_admin -d postgres -c "\l"</code>	List databases
<code>docker exec vikuna-postgres psql -U vikuna_admin -d kaala_dristi_db</code>	Open psql shell
<code>docker exec vikuna-postgres psql -U vikuna_admin -d postgres -c "SELECT version();" </code>	PG version
<code>docker exec vikuna-postgres psql ... -c "SELECT pg_size_pretty(pg_database_size('kaala_dristi_db'));"</code>	DB size

5.3 LLM VPS — Qwen3

Command	Description
<code>docker compose logs vikuna-llm --tail 30</code>	Check LLM logs
<code>docker compose restart vikuna-llm</code>	Restart LLM
<code>curl localhost:8080/health</code>	LLM health check

Test Inference (local)

```
curl -s http://localhost:8080/v1/chat/completions \
-H 'Content-Type: application/json' \
-H 'Authorization: Bearer vk-llm-d0efccfd15c0b8fd72214d0b9182032f84106b43' \
-d '{
  "messages": [
    {"role": "system", "content": "You are a trading assistant. /no_think"},
    {"role": "user", "content": "What is BANKNIFTY in one sentence?"}
  ],
  "max_tokens": 100,
  "temperature": 0.3
}' | jq '.choices[0].message.content'
```

Test Inference (public HTTPS)

```
curl -s https://llm.dristiq.io/v1/chat/completions \
-H 'Content-Type: application/json' \
-H 'Authorization: Bearer vk-llm-d0efccfd15c0b8fd72214d0b9182032f84106b43' \
-d '{"messages":[{"role":"user","content":"Hello /no_think"}],"max_tokens":50}' \
| jq '.choices[0].message.content'
```

6. DristiQ Product Context

6.1 Active FastAPI Endpoints (port 8101)

Endpoint	Purpose
/api/panchang/daily	Daily panchang card data (live)
/api/kd/*	KaalaDristi compute APIs

6.2 Key Compute Functions

Function	Status
compute_astro_daily_signals()	Live — full 2026 signals computed
compute_flow_intelligence()	RVOL guard pending post-M046
M046 volume quality migration	Written — pending execution after job schedules clear

6.3 Open Items

Item	Priority
Digitize May–Dec 2026 events into km_astro_calendar_2026 (3 images ready), re-run scoring	High
Execute M046 after job schedules clear	High
RVOL guard in compute_flow_intelligence() post-M046	High
AstroSignalBadge component on Dashboard	Medium
Update CLAUDE.md (migrations M063–M070 undocumented)	Medium
Intraday tab at /intraday/:indexId (EOD data, intraday-ready architecture)	Medium
Conflict engine — 7 cases implementation	Medium
3 missing tables: km_finastro_alerts, km_finastro_muhurta, km_astro_correlation	Low

7. LLM Service — Client Usage

Qwen3 4B running at llm.dristiq.io serves as a shared billable LLM inference service across Vikuna products and client projects.

7.1 Usage Guidelines

Rule	Detail
Always use /no_think	Add to system prompt for all structured output calls — suppresses chain-of-thought tokens
API Key per client	Issue separate API keys per client project for usage tracking (llama.cpp supports multiple keys via comma-separated list)
Temperature	0.3 for factual/structured outputs, 0.7 for creative/explanatory
Max tokens	100–300 for structured signals, 500+ for explanations
Concurrency	4 parallel slots available — suitable for low-concurrency client projects

7.2 N8N Integration

N8N workflows on the same VPS can call Qwen3 via the internal Docker network — zero network hop, zero SSL overhead:

```
# In N8N HTTP Request node:
URL: http://vikuna-llm:8080/v1/chat/completions
Method: POST
Auth: Bearer vk-llm-d0efccfd15c0b8fd72214d0b9182032f84106b43
```

7.3 Scale Path

Phase	Action
Now	Qwen3 4B on KVM2 (8 GB) — handles DristiQ + light client usage
Phase 2	Upgrade LLM VPS to KVM4 (16 GB) — more parallel slots, higher concurrency
Phase 3	Per-client API key billing — track tokens via llama.cpp logs

8. Upgrade Checklists

8.1 Main VPS — KVM2 → KVM4

Step	Action
1	Upgrade in Hostinger dashboard
2	Update .env: PG_SHARED_BUFFERS=4GB, PG_EFFECTIVE_CACHE=10GB
3	docker compose restart postgres
4	Set up pg-backup container with Google Drive
5	Configure SSL with Let's Encrypt via Nginx
6	Deploy node-backend (KI-Prime + FamilyKnows + VaNi)
7	Add swap (currently 0 — flagged risk on 8 GB during pipeline peaks)

8.2 LLM VPS — KVM2 → KVM4

Step	Action
1	Upgrade in Hostinger dashboard (Nov 2026 before expiry)
2	Update docker-compose: --threads 4 for vikuna-llm
3	Consider Qwen3 8B Q4_K_M (5 GB) for better quality at scale
4	Implement per-client API key billing

9. Cost Summary

Item	Monthly	Notes
Main VPS (KVM2)	₹799	Postgres + DristiQ + APIs
LLM VPS (KVM2)	₹799	Qwen3 + N8N
LLM inference	₹0	Self-hosted, CPU inference
Backups (Google Drive)	₹0	15 GB free tier (pending setup)
SSL (Let's Encrypt)	₹0	Via Traefik (LLM VPS) / pending Nginx (Main)
GROSS TOTAL	₹1,598	
Client LLM revenue (est.)	-₹500+	Billable per project
NET COST (est.)	₹1,098	After client offset

End of Document

Confidential | Vikuna Technologies | May 2026